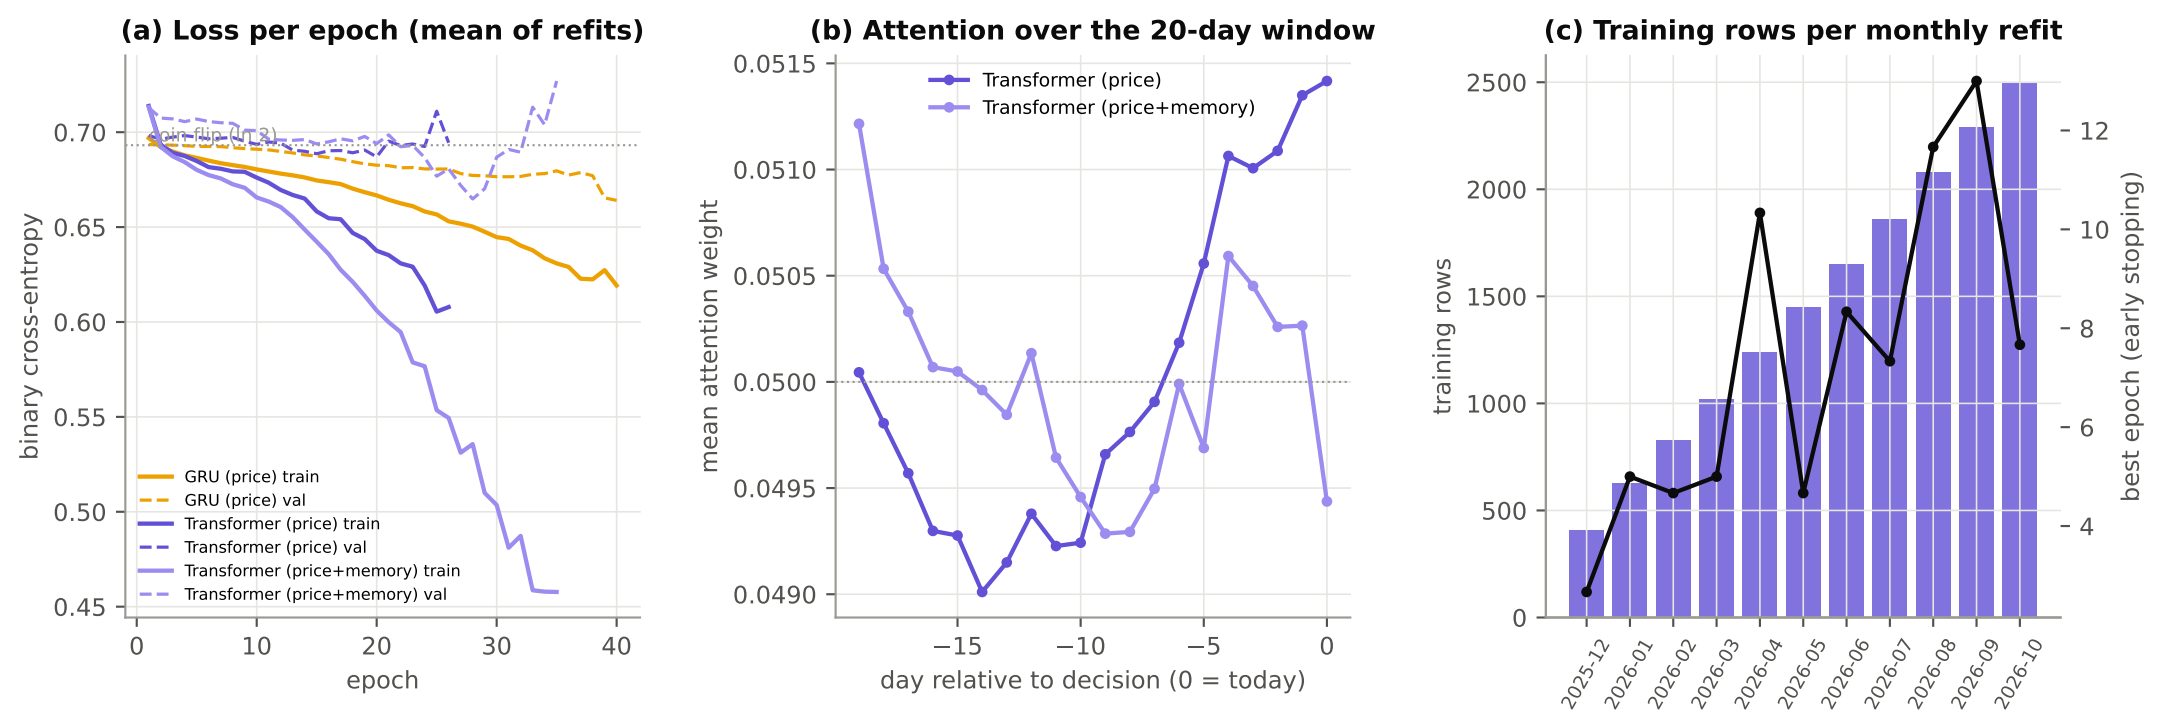
(a) Loss per epoch (mean of refits)
coin-flip (ln 2)
binary cross-entropy
epoch
GRU (price) train
GRU (price) val
Transformer (price) train
Transformer (price) val
Transformer (price+memory) train
Transformer (price+memory) val
(b) Attention over the 20-day window
Transformer (price)
Transformer (price+memory)
mean attention weight
day relative to decision (0 = today)
(c) Training rows per monthly refit
training rows
best epoch (early stopping)
2025-12
2026-01
2026-02
2026-03
2026-04
2026-05
2026-06
2026-07
2026-08
2026-09
2026-10